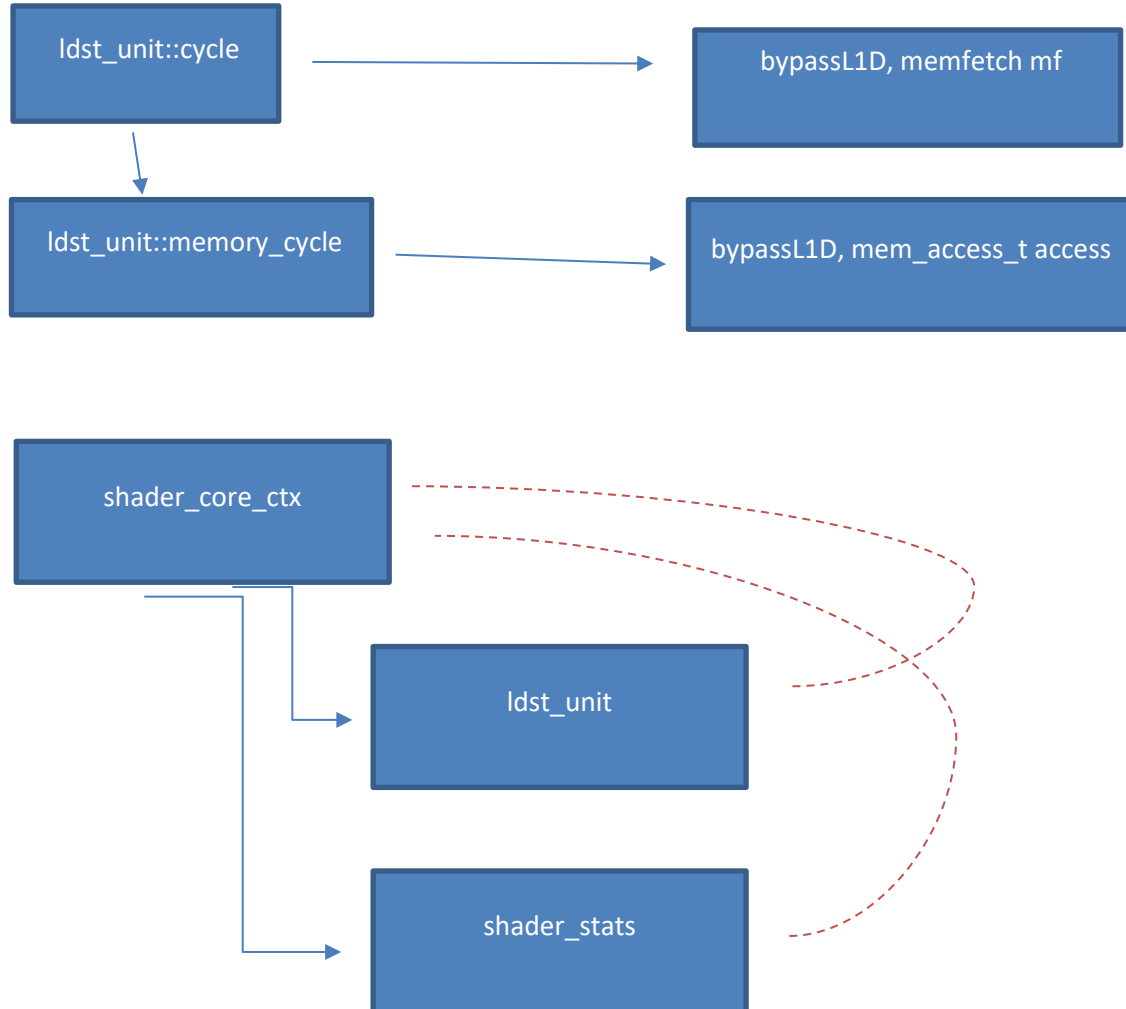


## ECE786: PROGRAMMING ASSIGNMENT THREE - A

### TASK: Load Bypass Mechanism

#### CODE ANALYSIS DIAGRAM



#### CODE ANALYSIS DISCUSSION

- The task requires to enforce L1 Data Cache bypass for loads with specific addresses. The load store part of gpgpu-sim is handled by the `ldst_unit` class that manages global, local, shared and texture memory accesses.
- The load instructions are issued for global and local memory. Investigating the methods of the `ldst_unit` class reveals two methods namely, `memory_cycle()` and `cycle()` that manage accesses to the L1 Data Cache within the simulator.
- The above diagram depicts the specific variables and class objects (`bypassL1D`, `access`, `mf`) that have been utilized to bypass the L1 Data Cache for the address range given in the assignment. This handles the first part of the assignment.
- It is also required to print the L1 bypasses for each kernel. Within a gpu each kernel is allocated to one SM. The SM is modelled by the `shader_core_ctx` class in the simulator.
- This SM class also contains its own instances of `ldst_unit` (`m_ldst_unit`) and `shader_stats` (`m_shader_stats`). These classes also have a pointer to the SM object for which they are instantiated as depicted in the above diagram,

- Based on the prior information, the counters for each kernel can be modelled in the following two ways:
  - First, a counter can be declared within the `shader_core_ctx` class which will be updated in the `memory_cycle` and `cycle` methods of `ldst_unit`. This requires building public methods to access the counter as the `shader_core_ctx` pointer is protected member in `ldst_unit`. In a similar way, `shader_stats` class can access this counter in its print method to print the bypasses for each kernel.
  - Second, a cumulative counter and a static counter declared in the `shader_stats` class. The cumulative counter is updated in the `ldst_unit` methods. The print method of `shader_stats` will access the cumulative counter and subtract it with static counter to get the bypasses for each kernel. The static counter stores the prior cumulative counter value.
- Because of the simplicity of the code, the second method has been utilized to implement the count mechanism for the bypasses.

### CODE SOLUTION SNIPPETS

```
for (unsigned i = 0; i < m_config->num_shader(); i++) {
    thread_icount_uarch += m_num_sim_insn[i];
    warp_icount_uarch += m_num_sim_winsn[i];
}
// *****
// print the L1 bypasses - Aditya
// *****

fprintf(fout, "***** L1 D$ Bypasses ***** : %d\n", this->count_cum_load_bypasses - sh
fprintf(fout, "***** L1 D$ Cum Bypasses ***** : %d\n", this->count_cum_load_bypasses);
shader_core_stats::count_load_bypasses_per_core= this->count_cum_load_bypasses;
fprintf(fout, "gpgpu_n_tot_thrd_icount = %lld\n", thread_icount_uarch);
fprintf(fout, "gpgpu_n_tot_w_icount = %lld\n", warp_icount_uarch);
```

```
class shader_core_stats : public shader_core_stats_pod {
public:

    // *****
    // counter to count load bypasses for each shader_core_ctx - Aditya
    // *****
    long count_cum_load_bypasses;

    // *****
    // counter to keep count of previous load bypasses for each shader core ctx - Aditya
    // *****
    static long count_load_bypasses_per_core;
    shader_core_stats(const shader_core_config *config) {
        // *****
        // Initialize counter to count load bypasses for each shader_core_ctx - Aditya
        // *****
        count_cum_load_bypasses = 0;

        m_config = config;
        shader_core_stats_pod *pod = reinterpret_cast<shader_core_stats_pod *>(
            this->shader_core_stats_pod_start);
        memset(pod, 0, sizeof(shader_core_stats_pod));
        shader_cycles = (unsigned long long *)calloc(config->num_shader(),
            sizeof(unsigned long long));
        m_num_sim_insn = (unsigned *)calloc(config->num_shader(), sizeof(unsigned));
```

```

// *****
// make and increment the counter for L1 Data Cache Load bypasses - Aditya
// *****
if(inst.is_load())
{
    new_addr_type addr = access.get_addr();
    if((addr >= 0xc0000000) && (addr <= 0xc00fffff)){
        bypassL1D = true;
        this->m_stats->count_cum_load_bypasses++;
    }
}

```

## RESULT

Cumulative Load Bypasses	Load Bypasses for each kernel
26	26
163	137
963	800
5884	4921
34108	28224
181692	147584
609201	427509
671956	62755
674391	2435
674406	15

Note: Sample count for one of the kernels

```

distro:
995, 721, 1039, 795, 806, 807, 935, 795, 827, 911, 1085, 1156, 923, 1505, 692, 798, 851,
825, 828, 922, 998, 1069, 848, 829, 283, 204, 304, 40, 241, 241, 272, 294,
***** L1 D$ Bypasses ***** : 15
***** L1 D$ Cum Bypasses ***** : 674406
gpgpu_n_tot_thrd_icount = 60452410
gpgpu_n_tot_w_icount = 1888513
gpgpu_n_stall_shd_mem = 1284090
gpgpu_n_mem_read_local = 0
gpgpu_n_mem_write_local = 0

```